

1                   

# Model Generation from Structured Light Scans

2                   Brandon Ha

3                   bha3@uci.edu

4                   June 12, 2020

## Abstract

5                   Converting two-dimensional scans to a three-dimensional model is no trivial task;  
6                   however, through the meticulous experimentation and application of various algo-  
7                   rithms and techniques, it is possible to generate satisfactory results as demonstrated  
8                   in this report. The report reinforces the idea that the quality of the result is heavily  
9                   dependent on every step of the process from collecting the initial scans and camera  
10                  calibration to surface reconstruction and additional clean-up processes.

11                

## 1 Project Overview

12                The problem in question revolves around the task of using a series of two-dimensional pictures  
13                to create something that is not only three-dimensional but also representative of the object in the  
14                given pictures. Therefore, the goal of the project is to take multiple scans of an object and create a  
15                three-dimensional object through a series of steps. More specifically, the majority of project entails  
16                the following:

- 17                • Reading structured light scans captured from two separate cameras  
18                • Decoding the structured light patterns from each scan  
19                • Generating meshes based on these decoded scans  
20                • Cleaning up and smoothing the meshes  
21                • Aligning and merging the generated meshes into one combined mesh  
22                • Performing Poisson surface reconstruction on the combined mesh

23                This is just a short list of the numerous intricacies required to produce a good quality model. Arguably,  
24                a significant reason as to how the result is able to look as accurate as it does is due to the time and effort  
25                dedicated to tweaking the countless variables and parameters throughout the many processes whether  
26                it be calibration, surface reconstruction, decoding, object mask computation, or mesh smoothing.

27                

## 2 Data

28                For this project, I opted to use a set of structured lights scans provided by Professor Charless Fowlkes  
29                due to the limited resources available at the time. Out of the four different objects he provided, I  
30                chose to create a model of the sculpture of a couple because I was interested in replicating the fine  
31                detail and surface texture of the object. Although the scans could be higher resolution, the resolution  
32                of the pictures at 1920 pixels by 1080 pixels prove to be adequate for our purposes.

33    **2.1 Calibration**

34    In order to calibrate the cameras, we were provided with 20 pictures from 20 different angles of a  
35    chessboard for both cameras. By varying the view of the chessboard, we can compute the intrinsic  
36    and extrinsic parameters of the cameras without having to create a complicated three-dimensional  
37    calibration setup.

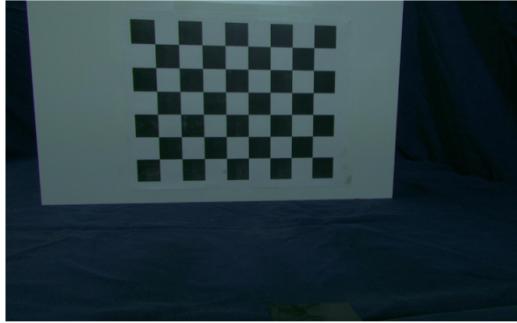


Figure 1 Left Camera - First View

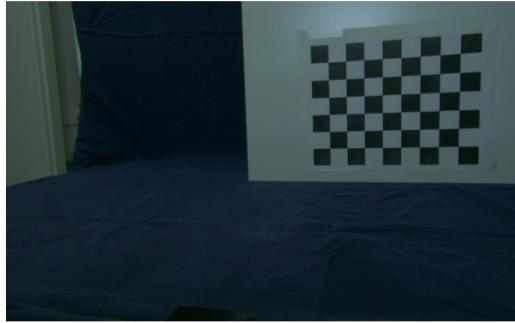


Figure 2 Right Camera - First View

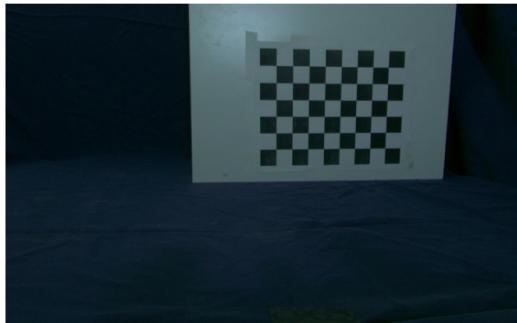


Figure 3 Left Camera - Second View

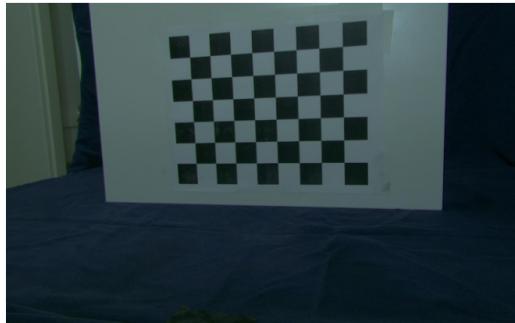


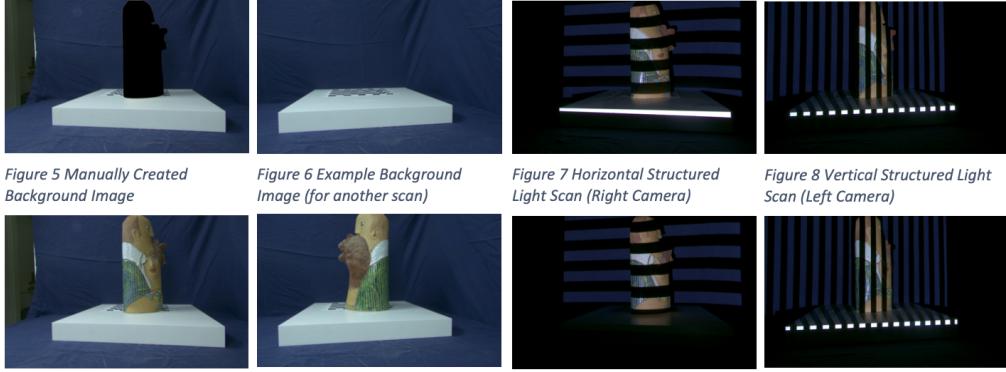
Figure 4 Right Camera - Second View

38

39    **2.2 Main Dataset**

40    The main dataset consisted of seven scans each composed of 84 images with the exception of one  
41    of the scans. Due to an error likely related to the image acquisition script used, the two images  
42    used for object mask computation, were missing. This meant that we could either disregard the scan  
43    completely or alternatively, manually create the missing image by filling where the object is with a  
44    color sufficiently different from any of the colors of the object. After alleviating this issue by masking  
45    the object completely with black pixels, we have all of the images necessary for the object masks and  
46    pixel decoding.

47    The seven scans are captured from different angles to maximize the coverage of the figure's surface.  
48    The scans can be broken down into two frames per camera where the scene is completely lit, one with  
49    the object and one without, and 40 frames per camera with projected structured light. By calculating  
50    the difference between the two frames, we can calculate which pixels the object occupies in the  
51    pictures for the scan. The 40 frames per camera can be further separated into projections of a binary  
52    gray code with 10 bits to project 10 different patterns, horizontally and vertically, and their respective  
53    inverses so that we can compare the corresponding frames and account for dark parts on the object.  
54    (Refer to Figures 5-12 on Page 3)



55

### 56 3 Algorithms

57 Although I have implemented many of the following algorithms I am about to discuss, I chose to  
 58 use the modules and functions provided by Professor Fowlkes because using the function provided  
 59 allowed me to focus on other aspects of the project such as configuring mesh smoothing, alignment,  
 60 and surface reconstruction.

61 **3.1 Camera Calibration**

62 The main algorithm used to calibrate the cameras remain mainly untouched from the provided  
 63 files. I loaded the calibration images and ensured that the calculated calibration parameters were  
 64 loaded correctly through the resulting pickle file. In terms of the methodology, I used OpenCV's  
 65 ***findChessboardCorners*** function to locate the chessboard corners and estimate the cameras' intrinsic  
 66 parameters, more specifically, the focal lengths and principal points of each camera. This is  
 67 possible because the geometry of the chessboard in the real world is known which leaves us with a  
 68 homogeneous linear system that OpenCV's ***calibrateCamera*** function solves.

69 After computing the intrinsic parameters, we need to estimate the extrinsic parameters of the cameras.  
 70 When you move the chessboard around, the intrinsic parameters stay the same, but the extrinsic  
 71 parameters (rotation and translation) for the two cameras change. With that said, the relative locations  
 72 of the two cameras stays constant, so we can use an arbitrary pair of calibration scans to estimate  
 73 the extrinsic parameters. For this project, I decided to use the first pair of calibration images and  
 74 ran the provided ***calibratePose*** function that essentially solves a least squares problem to minimize  
 75 the distance between the known two-dimensional points and the generated three-dimensional point  
 76 coordinates of the chessboard points.



77 *Figure 13 Calibration Visualization*

78 **3.2 Mesh Generation**

79 Generating high quality meshes is a critical step in the formation of a 3D model. Given that the mesh  
80 generation was explored in previous assignments, I was able to modify the provided code to handle  
81 color for each 3D point. Much of the mesh generation procedure is found in the **reconstruct** function  
82 while some aspects such as the internal function parameters are determined outside of the function.  
83 The next six subsections detail the main portions of the mesh generation process.

84 **3.2.1 Scan Decoding**

85 I used the provided function responsible for decoding the projected gray code for each frame and  
86 computing the mask indicating which pixels are likely to be inaccurately decoded. The function  
87 converts the images to grayscale and float value arrays after loading them in and then recovers a bit  
88 for every pixel that has a greater value in one image versus its inverse. Meanwhile, we maintain a  
89 binary array to mark difficult-to-decode pixels that have differences between pairs of scans above  
90 a given threshold. Finally, after going through all frames, I converted the resulting gray code to  
91 standard binary and used the binary code to calculate the decimal value.

92 **3.2.2 Object Mask Computation**

93 In terms of object mask computation, I heavily based my code on the provided mask generation  
94 implementation in the scan decoding function. The framework of the code can be fundamentally  
95 broken down into reading the background and foreground images, converting them to float values and  
96 grayscale, and lastly, thresholding the corresponding pairs of images. We can then use the generated  
97 masks with the already-computed decoding masks.

98 **3.2.3 Point Matching and Triangulation**

99 Once the cameras are calibrated and necessary masks are computed, I used NumPy's **intersect1d**  
100 function to find intersections between the left camera and right camera scans by inputting the  
101 respective masked codes. With the calculated indices, we map the intersecting points for the two  
102 cameras as well as the three-dimensional representation of said points with the provided triangulation  
103 function.

104 The **triangulate** function takes advantage of the fact that we know the intrinsic camera properties and  
105 the 2D coordinates of each matching point. This becomes a least squares problem which can easily  
106 be solved with a call to NumPy's least squares solver. Averaging the left camera and right camera  
107 results gives us a good estimate of the three-dimensional point coordinates.

108 After the 2D and 3D points are computed, I made sure to pull the color values from the original  
109 foreground scan for later use when writing to the mesh files (ending in .ply).

110 **3.2.4 Boundary Box Pruning**

111 Depending on the chosen thresholds for object mask computation and decoding, you will get varying  
112 amounts of noise and point density. To remove clear outliers due to low thresholds, boundary box  
113 pruning is used. Extensive trial and error is needed for this process. The general approach requires  
114 2D graphs of the cameras in relation to the triangulated points. Boundary box pruning is a relatively  
115 hands-on strategy in which the user manually selects hard limits on each coordinate plane to remove  
116 extrema. My approach involves making a mask that marks points that violate the user-provided limits  
117 and removing the offending points.

118 **3.2.5 Delaunay Triangulation**

119 To form the triangles for the mesh, I used SciPy's Delaunay triangulation function due to its simplicity  
120 and effectiveness. Delaunay triangulation is a very diverse algorithm for computing triangulation and  
121 is known for its ability to form relatively even triangles. This triangulation evenness is a result of  
122 the Delaunay criterion which follows the empty circumcircle criterion. According to the circumcircle  
123 criterion, no point outside of a triangle's vertices may be located within the circumcircle of its

124 corresponding triangle.<sup>1</sup> Thus, triangles with larger angles are favored over those with smaller  
125 angles. This serves as a good base for the mesh, but the algorithm can still form long triangles.

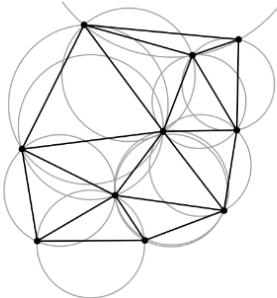


Figure 14 Delaunay Triangulation  
with Circumcircles

126

2

### 127 3.2.6 Triangle Pruning

128 Triangulation can result in undesirably long triangles which can reduce mesh quality. As a result, I  
129 used code from a previous assignment to prune these triangles. I made a mask which indicated which  
130 triangles with edges shorter than a given threshold. Then, I only kept those triangles to process. In  
131 doing so, it leaves points that are not referred by any triangles.

### 132 3.2.7 Removing Unreferenced Points

133 Because triangle pruning can leave many unreferenced points, it is important to remove points that  
134 fall in this criteria. Following Professor Fowlkes notes, I was able to remap the remaining vertices  
135 listed in the array of triangles. The code in question can be seen below.

```
136 # Remapping triangles
137 tokeep = np.unique(tri)
138 remap = np.full(pts3.shape[1], -1)
139 remap[tokeep] = np.arange(0, tokeep.shape[0])
140 tri = remap[tri]
141
142 # Remove loose points
143 pts3 = pts3[:,tokeep]
144 ptsColor = ptsColor[:,tokeep]
145 pts2L = pts2L[:,tokeep]
146 pts2R = pts2R[:,tokeep]
```

## 147 3.3 Mesh Smoothing

148 While the mesh is fully generated at this point, the quality tends to be quite low without some form of  
149 smoothing. So, I created a function to smooth each generated mesh. The general idea is to record the  
150 neighbors of every point. With this information stored in a dictionary, I was able to index into the  
151 3D point array and calculate the average coordinate for any given points' neighbors. This greatly  
152 improves the smoothness of the mesh and minimizes any remaining floating point clusters.

---

<sup>1</sup><https://www.mathworks.com/help/matlab/math/delaunay-triangulation.html>

<sup>2</sup>[https://en.wikipedia.org/wiki/Delaunay\\_triangulation](https://en.wikipedia.org/wiki/Delaunay_triangulation)

153 **4 Results**

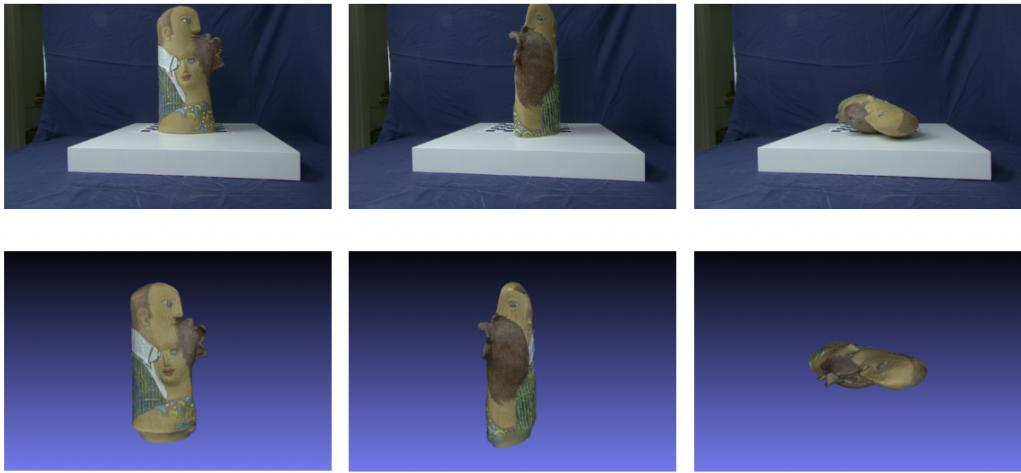


Figure 15 Original Scans versus  
Final Mesh

154

155 **4.1 Final Model**

156 Looking at the final result and comparing it side-by-side to the original object, we can clearly see that  
157 the generated model captures the basic geometry and texture of the original object. There are spots  
158 where the model misrepresents portions of the figure, mainly the bottom and the top of the object.  
159 The provided scans did cover most of the surface of object, but the two spots where the model has the  
160 most difficulty are the spots where the scans fail to capture. Again, near the bottom of the model, you  
161 can see that some of the patterning in the woman's clothing is somewhat blurry. With that being said,  
162 the model proves to be quite detailed and even displays dimensionality throughout as shown by the  
163 height of woman's nose and the green lines in the man's jacket.

164 **4.2 Effects of Processing**

165 Through my research and experimentation, I learned that reconstruction quality is heavily reliant on  
166 many different aspects of the process.

167 **4.2.1 Boundary Box and Triangle Pruning**



Figure 16 Original Mesh

Figure 17 Mesh without Boundary Box Pruning

Figure 18 Mesh without Triangle Pruning

168

169 Although boundary box pruning and triangle pruning act as simple forms of processing, the effect  
170 they have on the mesh is undeniably remarkable. Boundary box pruning removes many of the points  
171 that get past the different masks. While you could change the object mask and decoding thresholds  
172 to minimize this kind of noise, it can compromise the point density of the object we are trying to  
173 resolve. Triangle pruning removes the noise that is around the object elegantly as seen in Figure 18.

174 **4.2.2 Mesh Smoothing**



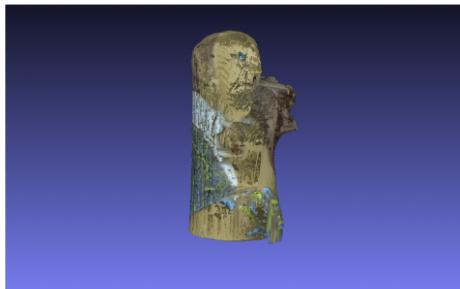
175 *Figure 19 Mesh without Smoothing*



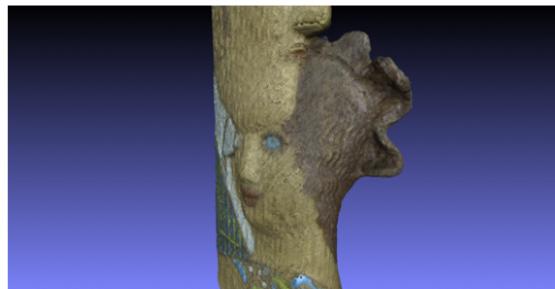
*Figure 20 Aligned Model without Smoothing*

176 Smoothing also plays an important role in the final mesh. Comparing the original mesh with the mesh  
177 without smoothing (Figures 16 and 19), you may be able to notice the graininess of the unprocessed  
178 mesh. This effect is further exacerbated once the meshes are merged and surface reconstruction is  
179 performed.

180 **4.2.3 Surface Reconstruction**



181 *Figure 21 Merged Mesh without Surface Reconstruction*



*Figure 22 Poisson Surface Reconstruction with Reconstruction Depth = 12*

182 Poisson surface reconstruction brings together the merged meshes and creates a mesh that is smoother  
183 and more refined by solving an implicit function in which the normal vectors on the surface of the  
184 object are represented by a gradient. Figure 21 shows that without surface reconstruction, the model  
185 is not closed and has texture issues because the different meshes are visible in different patches on  
186 the surface of the object. MeshLab's Poisson surface reconstruction handles this issue by taking  
187 the average color for each corresponding point in each mesh. Additionally, I went through different  
188 values for reconstruction depth which profoundly changed the smoothness of the model's surface.  
189 At a reconstruction depth of 12, there was too much artificial surface texture which is why I kept  
190 experimenting until I found an accurate result at a depth of 9. What we are left with is a relatively  
191 high-quality 3D model.

192 **5 Assessment and Evaluation**

193 **5.1 Limitations of Approach**

194 The approach I ended up using proved to lack much configurability despite the inclusion of the  
195 thresholds and parameters discussed in earlier sections. The final mesh still has blurring issues and  
196 does not recognize that the bottom of the figure does not extend as far as the model suggests. Without  
197 more user-provided input or more sophisticated surface reconstruction methods, these issues will  
198 continue to bottleneck the system's ability to create high-quality models.

199 **5.2 Limitations of Data**

200 The dataset provided for this figure is lacking in that the scans do not completely cover the full  
201 surface of the original object. If the appropriate scans were provided, then the model would be  
202 more geometrically correct. Resolution seem to sufficient as detail is retained even after the mesh  
203 smoothing.

204 **5.3 Final Thoughts**

205 Given the resources available and significant amount of time spent on this project, I believe that the  
206 final result is more than satisfactory. The dataset provided did leave some work for the algorithms  
207 to improvise and as such, the model was unable to completely render the bottom and part of the  
208 top of the object. With that being said, the result does capture the overall essence of the original  
209 figure. The task would be considered as a relatively easy task for a computer algorithm to solve when  
210 you consider the vast range of problems that a computer algorithm *can* solve. If I had more time, I  
211 would likely use another algorithm to smooth the meshes and explore different surface reconstruction  
212 methods. It became evident that my current implementation of the mesh smoothing lacks control  
213 other than the amount of iterations it should perform. Consequently, I view it as one of the weakest  
214 links in the pipeline.

215 **6 Appendix**

216 The following section serves as a quick guide for the code and process used in the project. If you  
217 would like to see more of the code, please refer to the source code. The relevant functions have  
218 comments that should clarify any confusions.

219 **6.1 Calibration**

220 Calibration is mostly unchanged from the provided code in **calibrate.py**. Some code from Assignment  
221 3 is used to estimate the extrinsic parameters as shown here:

```
Calibration

[5]: if not os.path.exists("calibration.pickle"):
    import calibrate
    with open('calibration.pickle', "rb") as calib_file:
        calib_data = pickle.load(calib_file)
    f = (calib_data['fx'] + calib_data['fy']) / 2
    c = np.array([calib_data['cx']], [calib_data['cy']]))

# load in the left and right images and find the coordinates of
# the chessboard corners using OpenCV
imgL = cv2.imread('../calib_jpg_l/frame_C_0_01.jpg')
ret, cornersL = cv2.findChessboardCorners(imgL, (8,6), None)
pts2L = cornersL.squeeze()

imgR = cv2.imread('../calib_jpg_r/frame_C_1_01.jpg')
ret, cornersR = cv2.findChessboardCorners(imgR, (8,6), None)
pts2R = cornersR.squeeze().T

# generate the known 3D point coordinates of points on the checkerboard in cm
pts3 = np.zeros((3,6*8))
yy,xx = np.meshgrid(np.arange(8),np.arange(6))
pts3[0,:] = 2.8*xx.reshape(1,-1)
pts3[1,:] = 2.8*yy.reshape(1,-1)

# Camera calibration (extrinsic parameters)
paramsL = np.array([0,0,0,0.5,-0.5])
paramsR = np.array([0,0,0,0.5,-0.5])

camL = camutils.Camera(f=f,c=c,R=camutils.makerotation(0,0,0),t=np.zeros((3,1)))
camR = camutils.Camera(f=f,c=c,R=camutils.makerotation(0,0,0),t=np.zeros((3,1)))

camL = camutils.calibratePose(pts3,pts2L,camL,paramsL)
camR = camutils.calibratePose(pts3,pts2R,camR,paramsR)

Estimated camera intrinsic parameter matrix K
[[1.40530791e+03  0.00000000e+00  9.6167369e+02]
 [0.00000000e+00  1.40389402e+03  5.90915957e+02]
 [0.00000000e+00  0.00000000e+00  1.00000000e+00]]
Estimated radial distortion coefficients
[[-5.23053889e-03  8.05964692e-02  2.01737110e-05 -3.90735917e-03
 -1.08096617e-01]]
Individual intrinsic parameters
fx = 1405.3079107410456
fy = 1403.8940216103824
cx = 962.167369266617
cy = 596.9159568140669
```

223 **6.2 Mesh Generation**

224 Mesh generation is a modified version of the my mesh generation code in Assignment 4 as well as  
 225 the provided **reconstruct** function in **camutils.py**.

226

**Mesh Creation Function**

```
[4]: def mesh(scanFolder, camL, camR, params, output):
    """
    Mesh generation based on given scans, two cameras,
    and various parameters that allow for tweaking. Writes
    mesh to .ply file.

    Parameters
    -----
    scanFolder : str
        folder for where the images are stored
    camL, camR : Camera
        camera parameters
    params : list of ints
        wrapper for mesh generation parameters
    output : str
        output filename

    mesh_num : int
        scan number
    ...
    print("Current folder:", scanfolder)
    print("Parameters for scan:", params)
    print("Output filename:", output)
```

226

**Modified reconstruct Function for Handling Color**

227

```
[2]: def reconstruct(scanFolder, mask_thres, decode_thres, camL, camR):
    """
    Reconstruction based on triangulating matched pairs of points
    between to view which have been encoded with a 20bit gray code
    for colored images.

    Parameters
    -----
    scanFolder : str
        folder for where the images are stored
    mask_thres : int
        object mask threshold
    decode_thres : float
        decodability threshold
    camL, camR : Camera
        camera parameters
    Returns
    -----
    pts2L, pts2R : 2D numpy.array (dtype=float)
    pts3 : 2D numpy.array (dtype=float)
    ptsColor : 2D numpy.array (3xN array where N = number of points)
    ...
    """


```

227

228

```
[6]: scansFolder = "./couple"
scanCount = len(glob.glob(os.path.join(scansFolder, "*")))

# Parameters for mesh generation
mask_thresholds = [0.005, 0.001, 0.001, 0.005, 0.0025, 0.001, 0.005, 0.005]
decode_thresholds = [0.005] * scanCount
trithresh
smooth_cnt = 2
boxLim = np.array([[-15, 24, 0.5, 22, -27.5, -15],
                  [-12, 22, 4.5, 17.5, -27, -10],
                  [-15, 23, 4.5, 22, -28, -10],
                  [-12, 23, -100, 100, -22, -10],
                  [-12, 24, 5, 1, 20, -38, -10],
                  [13, 5, 21, -10, 22, -39, 1]])]

for i in range(scanCount):
    params = [mask_thresholds[i], decode_thresholds[i], boxLim[i], trithresh, smooth_cnt]
    scanFolder = f"{scansFolder}/{i}"
    output_fname = f'{scansFolder}/{i}.ply'
    mesh(scanFolder, camL, camR, params, output_fname)
```

229 **6.3 Mesh Smoothing**

230 I wrote the mesh smoothing process which records each points' neighbors and changes the coordinates  
 231 of each point to the average location between its neighbors.

```

Mesh Smoothing Function

(3): def build_neighbor_dict(tri):
    # Auxiliary function to build dict of points and corresponding neighbors
    result = defaultdict(set)

    for triangle in tri:
        for pt_idx in triangle:
            current = {point for point in triangle if point != pt_idx}
            result[pt_idx] = result[pt_idx].union(current)

    return result

def smooth_mesh(pts3,tri,iterCount):
    """
    Iterative mesh smoothing via averaging l-ring neighbors
    Parameters
    -----
    pts3 : 2D numpy.array (dtype=float)
        tri : array of triangles vertices
    iterCount : int
        count of smoothing iterations
    Returns
    -----
    pts3 : 2D numpy.array (dtype=float)
    """
    neighbor_dict = build_neighbor_dict(tri)
    pts3_t = pts3.T

    for _ in range(iterCount):
        for idx, point in enumerate(pts3_t):
            neighbors = list(neighbor_dict[idx])
            sub = np.array([pts3_t[i,:] for x in neighbors])
            sub = np.average(sub, axis=0)
            pts3_t[idx,:] = sub

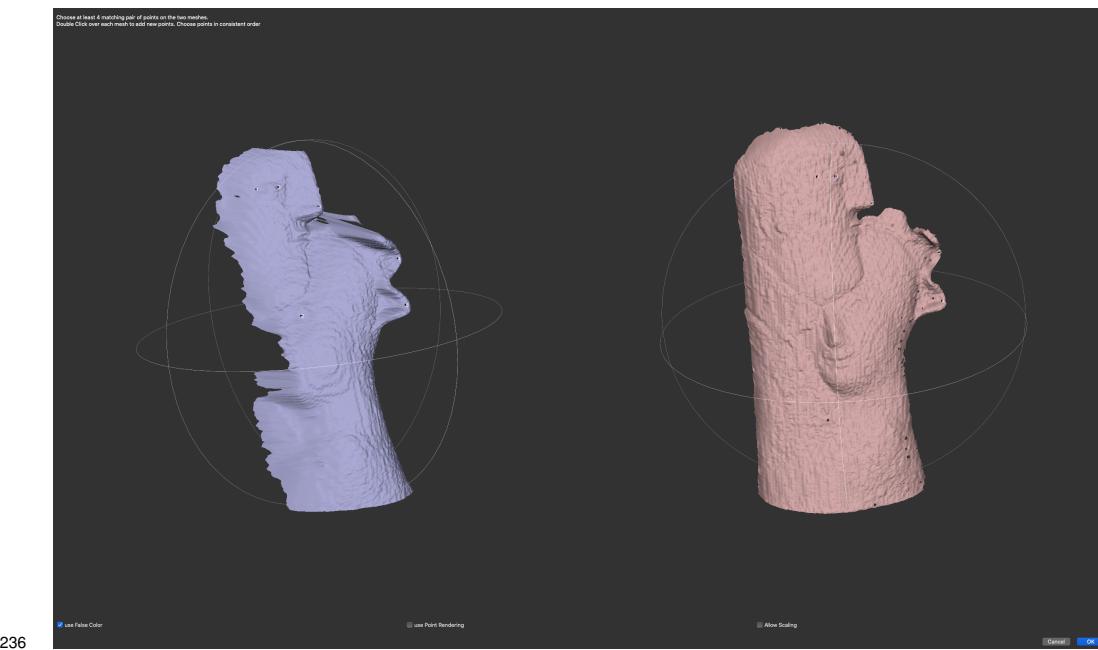
    return pts3_t.T

```

232

## 233 6.4 Mesh Alignment

234 I used MeshLab's mesh alignment tool to align every mesh. Then, I flattened the layers to create one  
 235 merged mesh.



## 237 6.5 Poisson Surface Reconstruction

238 While I initially wanted to use PyPoisson to perform surface reconstruction, it failed to compile in  
 239 my environment and lead to run-time issues when forced to compile with up-to-date dependencies.  
 240 So, I decided to resort to MeshLab's Poisson Surface Reconstruction method.